

Misc

- **Exercise 2 updated with Part III.**
 - Due on next Tuesday 12:30pm.
- **Project 2 (Suggestion)**
 - Write a small test for each call.
 - Start from file system calls with simple Linux translation.
 - Or start from Exec()
- **Midterm (May 7)**
 - Close book. Bring 3 pages of double-sided notes.
 - Next Monday: Project 2 or review?

More on Address Translation



CS170 Fall 2015. T. Yang

Based on Slides from John Kubiawicz

<http://cs162.eecs.Berkeley.edu>

Implementation Options for Page Table

- **Page sharing among process**
- **What can page table entries be utilized?**
- **Page table implementation**
 - One-level page table
 - Multi-level paging
 - Inverted page tables

Shared Pages through Paging

- **Shared code**

- One copy of read-only code shared among processes (i.e., text editors, compilers, window systems).
- Shared code must appear in same location in the logical address space of all processes

- **Private code and data**

- Each process keeps a separate copy of the code and data
- The pages for the private code and data can appear anywhere in the logical address space

Shared Pages Example

ed 1
ed 2
ed 3
data 1

process P_1

3
4
6
1

page table
for P_1

ed 1
ed 2
ed 3
data 2

process P_2

3
4
6
7

page table
for P_2

ed 1
ed 2
ed 3
data 3

process P_3

3
4
6
2

page table
for P_3

0	
1	data 1
2	data 3
3	ed 1
4	ed 2
5	
6	ed 3
7	data 2
8	
9	
10	
11	

Example

Virtual Address
(Process A):

Virtual Page #	Offset
----------------	--------

PageTablePtrA

page #0	V,R
page #1	V,R
page #2	V,R,W
page #3	V,R,W
page #4	N
page #5	V,R,W

PageTablePtrB

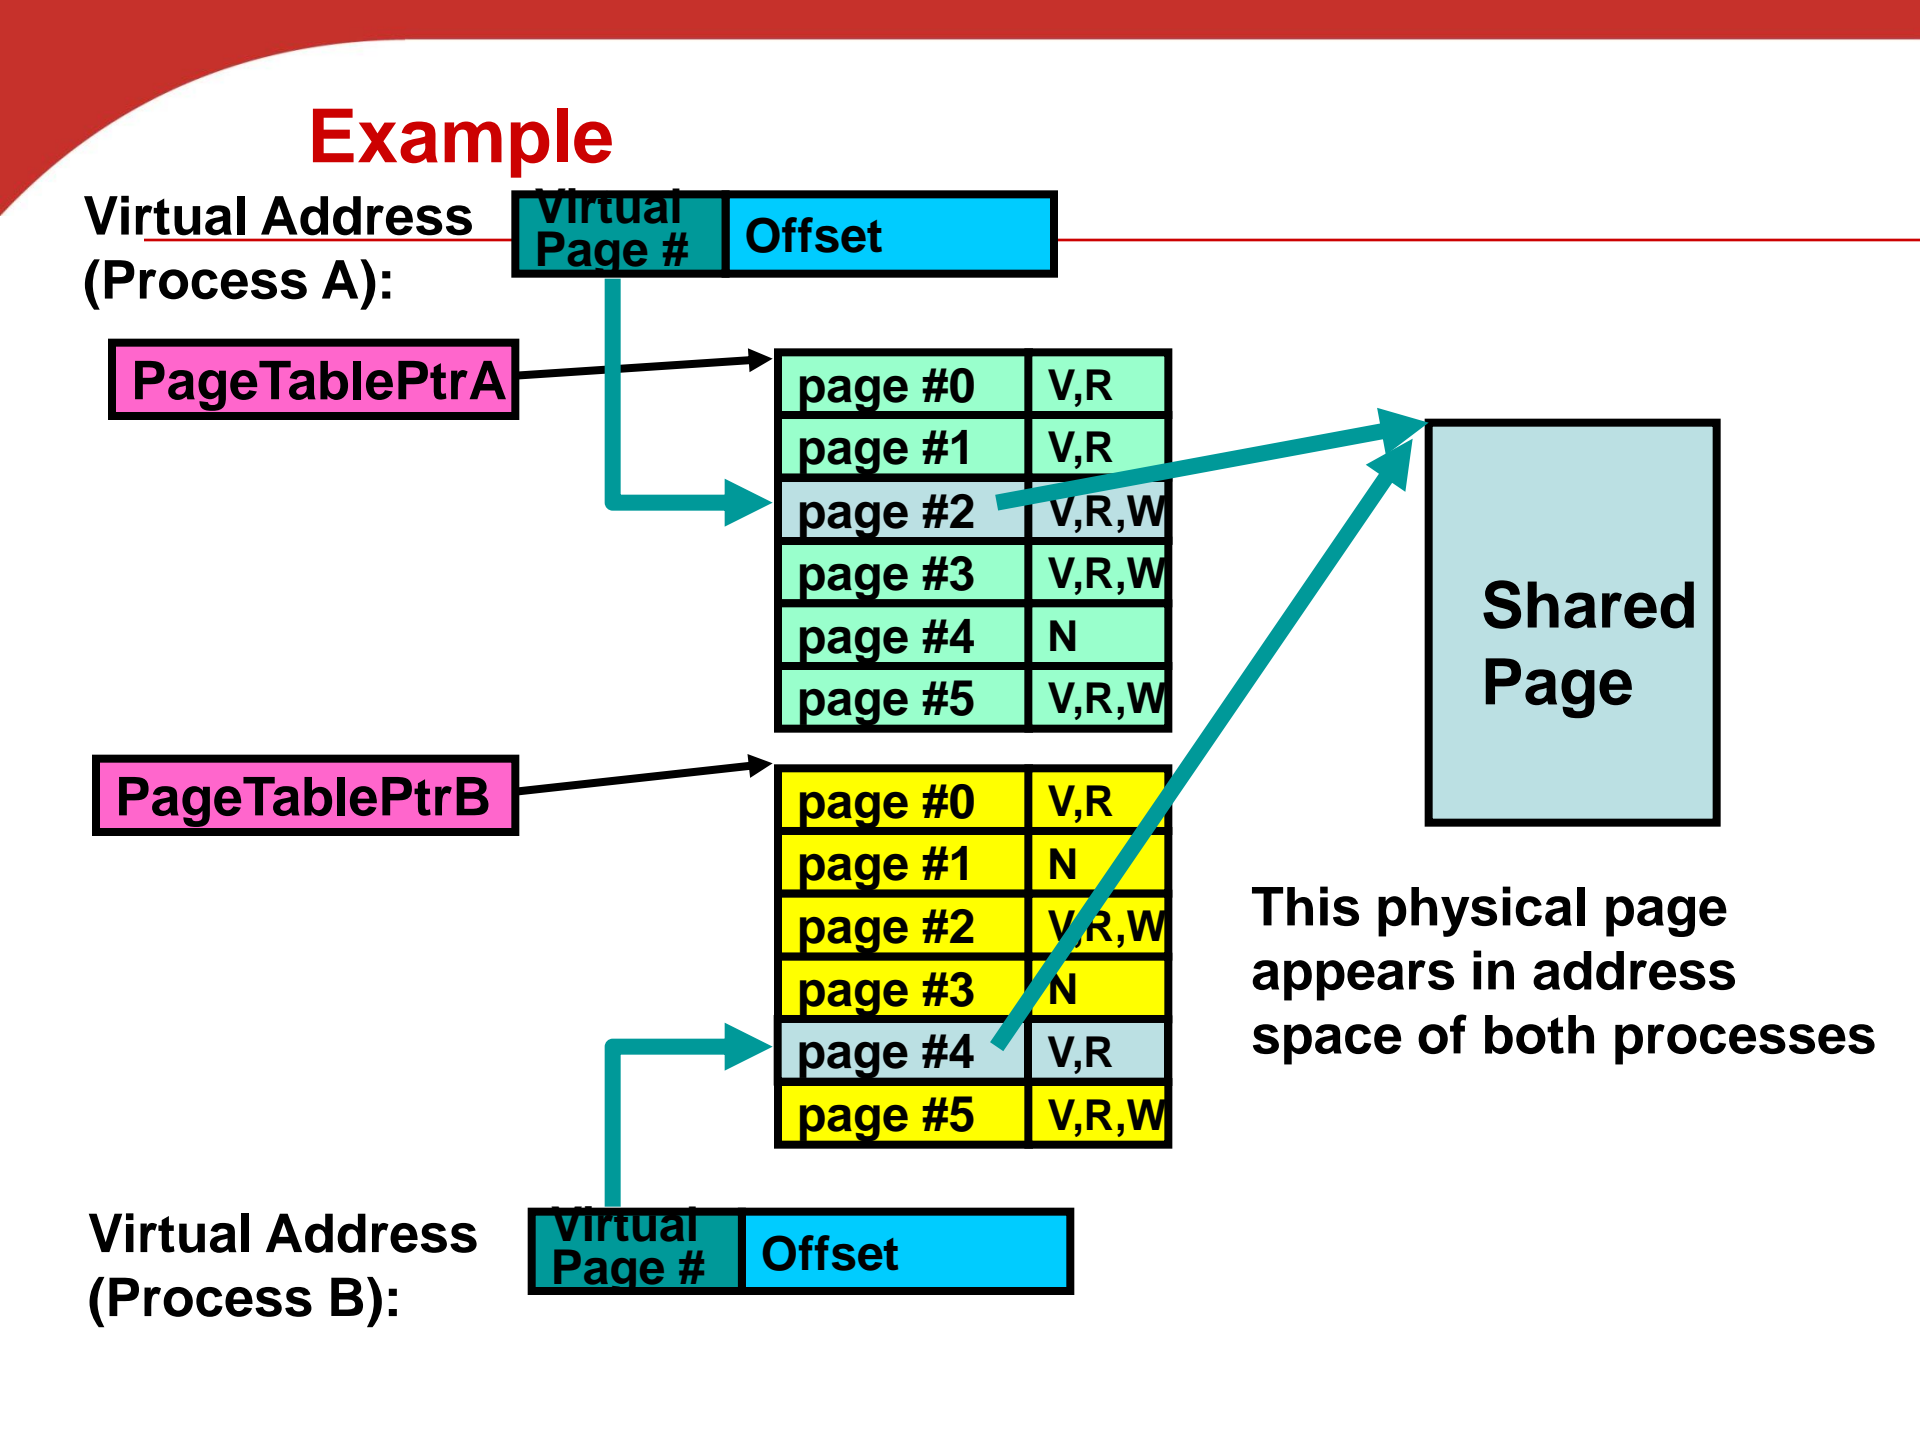
page #0	V,R
page #1	N
page #2	V,R,W
page #3	N
page #4	V,R
page #5	V,R,W

Virtual Address
(Process B):

Virtual Page #	Offset
----------------	--------

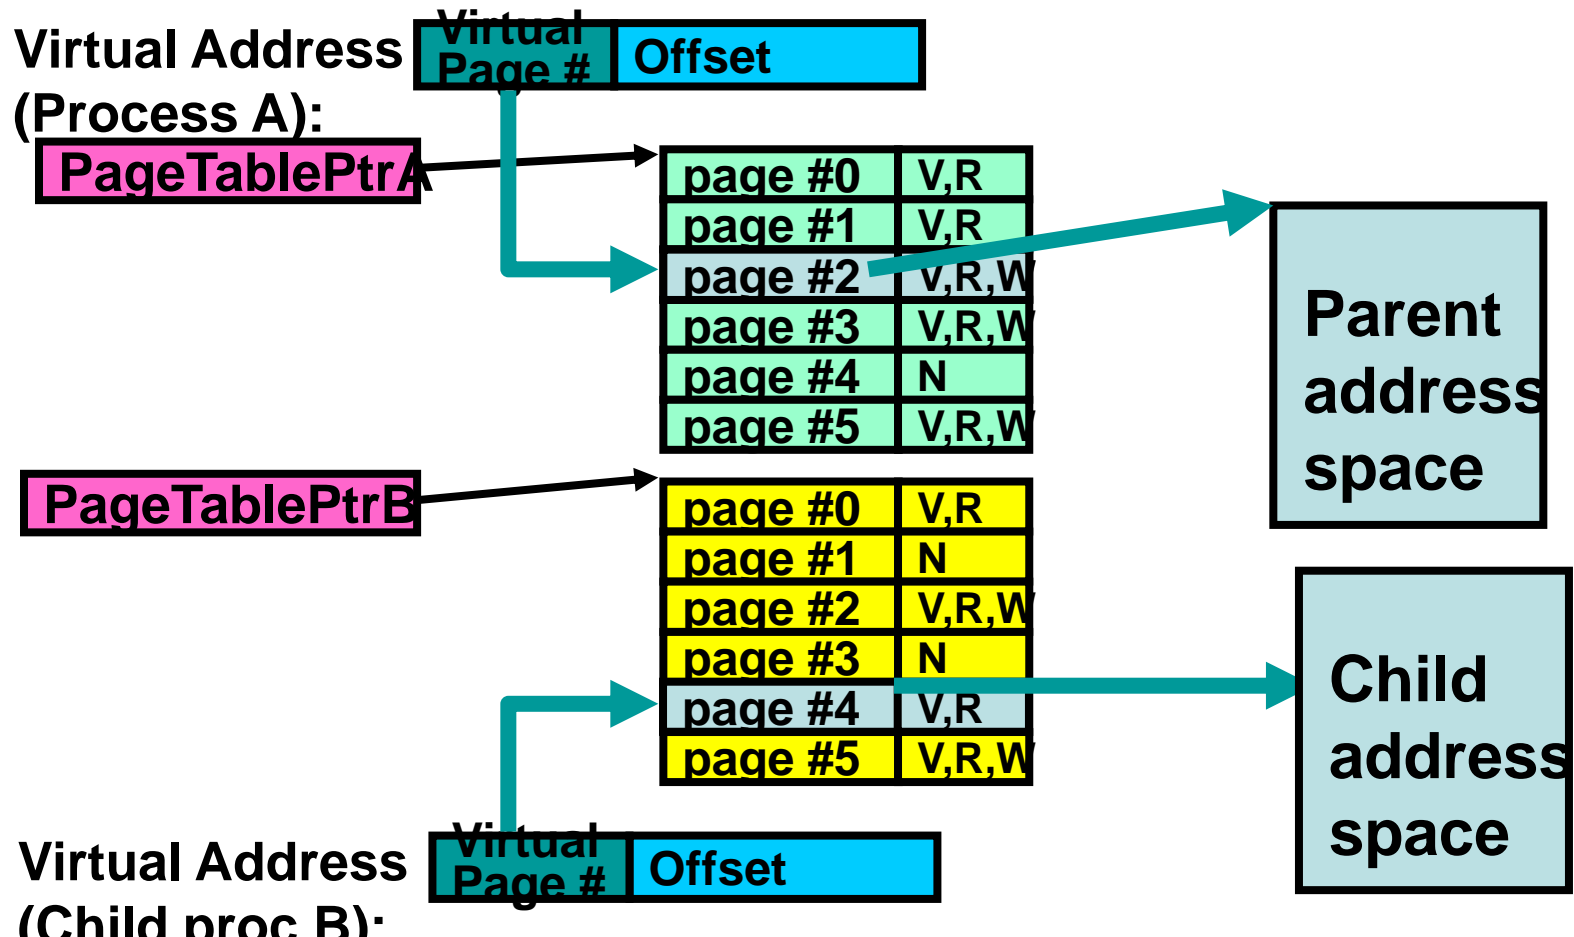
Shared
Page

This physical page
appears in address
space of both processes

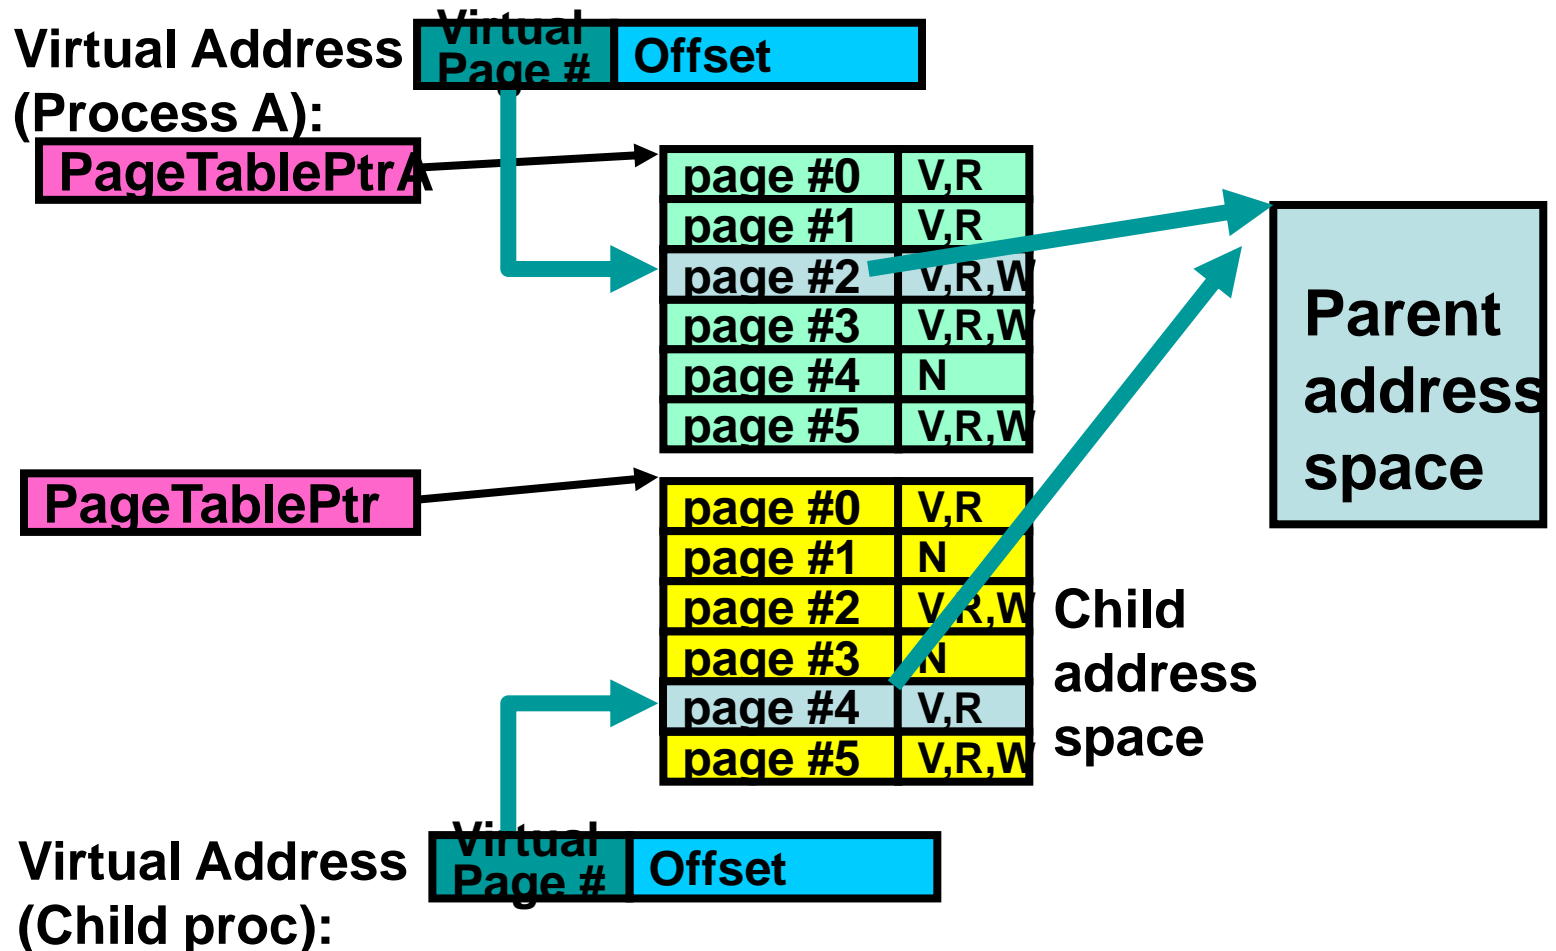


Optimization of Unix System Call Fork()

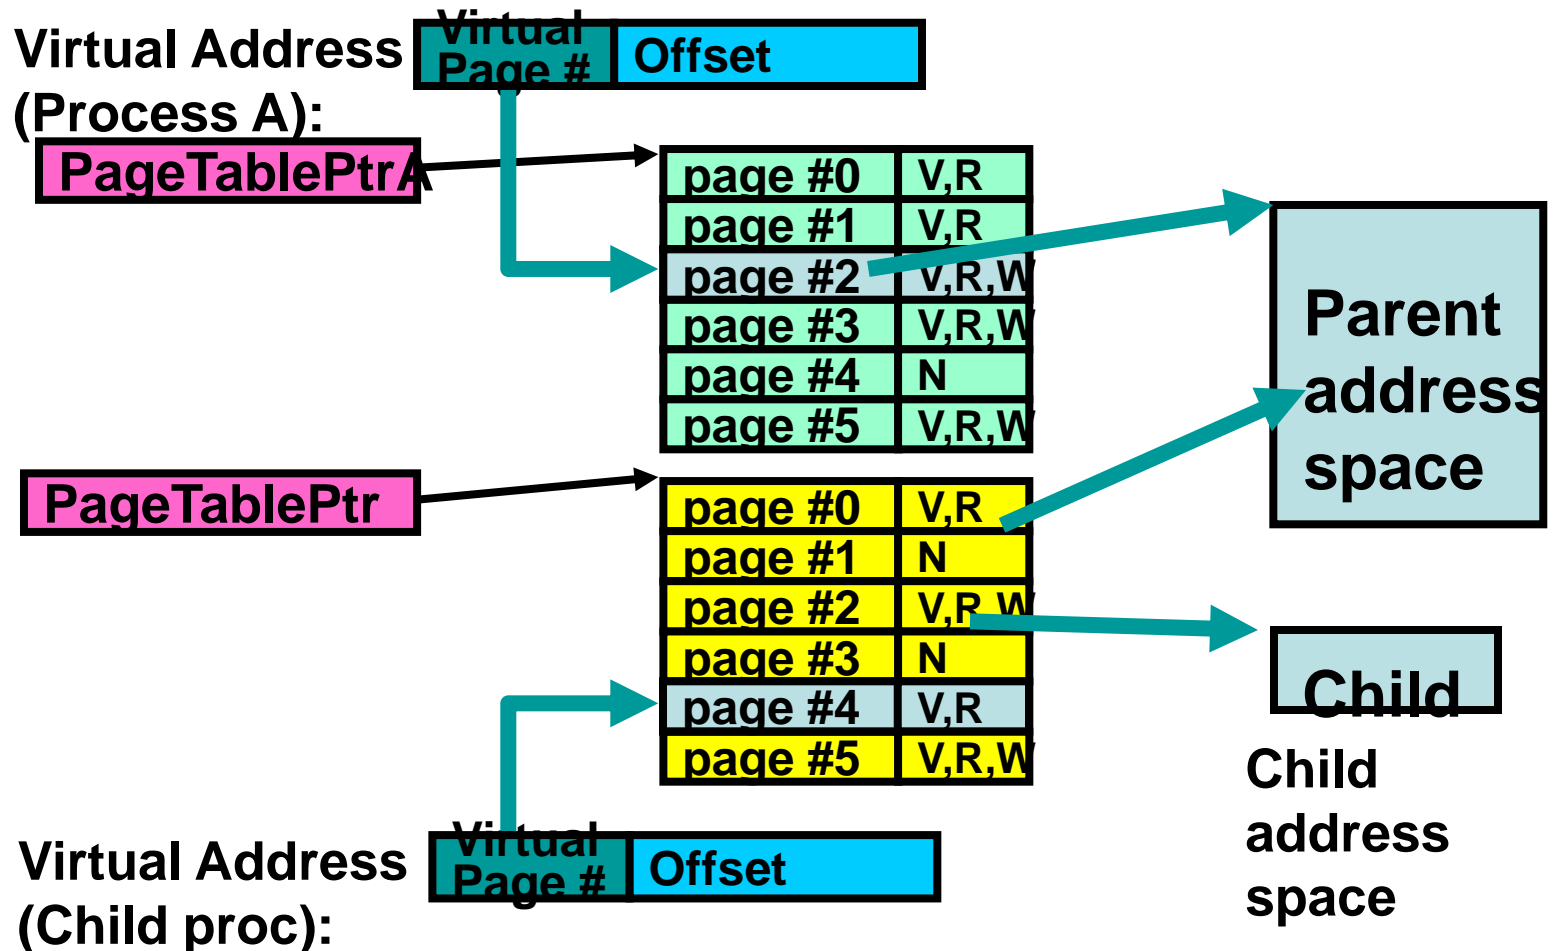
- A child process copies address space of parent.
 - Most of time it is wasted as the child performs `exec()`.
 - Can we avoid doing copying on a `fork()`?



Unix fork() optimization

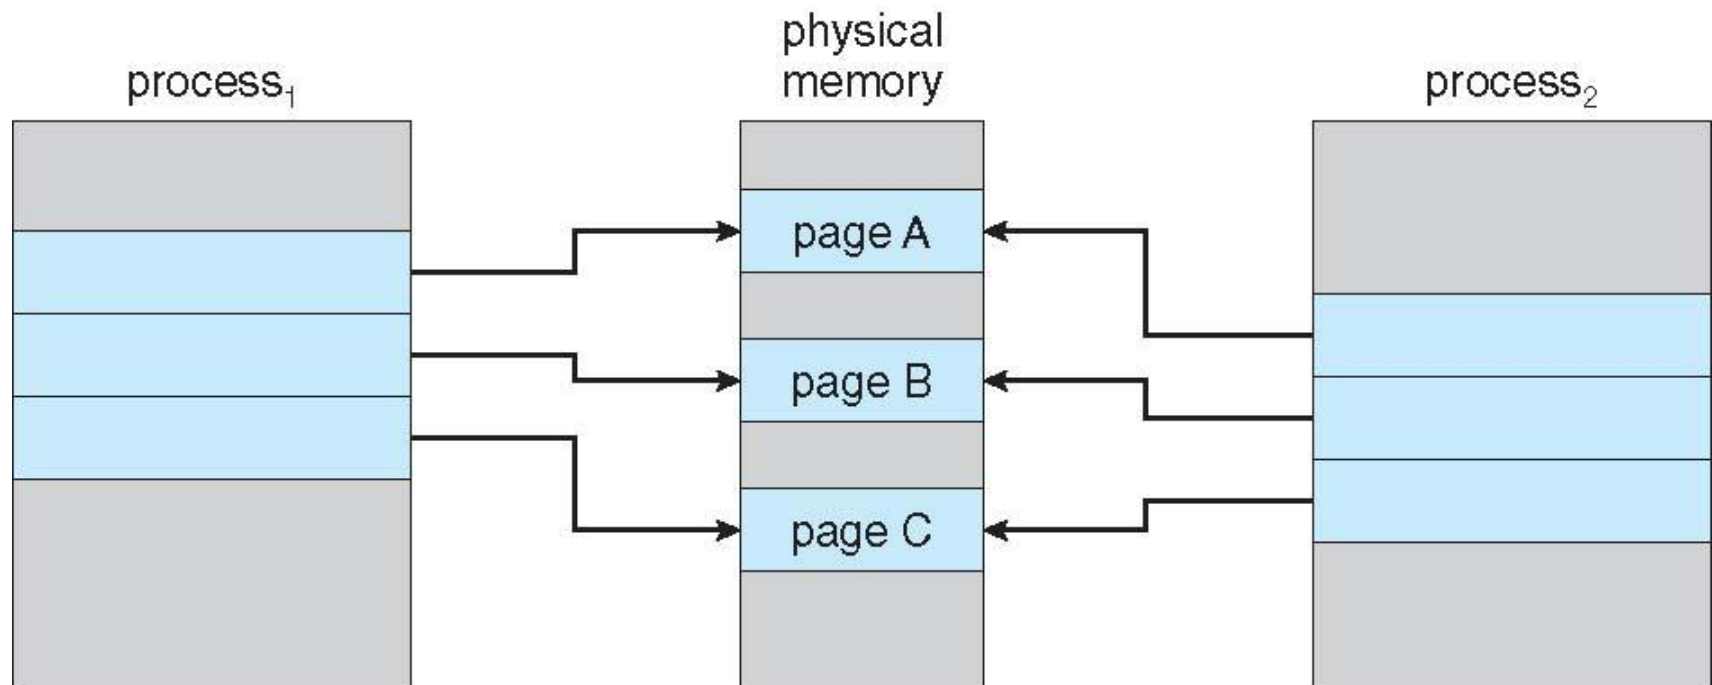


Unix fork() optimization




Copy-on-Write: Lazy copy during process creation

- **COW** allows both parent and child processes to initially *share* the same pages in memory.
- A shared page is duplicated only when modified
- COW allows more efficient process creation as only modified pages are copied



Copy on Write: After Process 1 Modifies Page C



How to memorize a page is shared?

When to detect the need for duplication?

Need a page table entry bit

Page table entry

Physical page number

--	--	--	--

More examples of utilizing page table entries

Page table entry	Physical page number			
------------------	----------------------	--	--	--

- **How do we use the PTE?**
 - Invalid PTE can imply different things:
 - Region of address space is actually invalid or
 - Page/directory is just somewhere else than memory
 - Validity checked first
 - OS can use other bits for location info
- **Usage Example: Copy on Write**
 - Indicate a page is shared with a parent
- **Usage Example: Demand Paging**
 - Keep only active pages in memory
 - Place others on disk and mark their PTEs invalid

Example: Intel x86 architecture PTE

- Address format (10, 10, 12-bit offset)
- Intermediate page tables called “Directories”

Page Frame Number (Physical Page Number)	Free (OS)	0	L	D	A	PCD	PWT	U	W	P
31-12	11-9	8	7	6	5	4	3	2	1	0

P: Present (same as “valid” bit in other architectures)

W: Writeable

U: User accessible

PWT: Page write transparent: external cache write-through

PCD: Page cache disabled (page cannot be cached)

A: Accessed: page has been accessed recently

D: Dirty (PTE only): page has been modified recently

L: L=1 \Rightarrow 4MB page (directory only).


Bottom 22 bits of virtual address serve as offset

More examples of utilizing page table entries

- **Usage Example: Zero Fill On Demand**
 - Security and performance advantages
 - New pages carry no information
 - Give new pages to a process initially with PTEs marked as invalid.
 - During access time, page fault → physical frames are allocated and filled with zeros
 - Often, OS creates zeroed pages in background
- **Can a process modify its own translation tables?**
 - **NO!**
 - If it could, could get access to all of physical memory
 - Has to be restricted

Implementation Options for Page Table

- **Page sharing among process**
- **What can page table entries be utilized?**
- **Page table implementation**
 - One-level page table
 - Multi-level paging
 - Inverted page tables




One-Level Page Table

- What is the maximum size of logical space?

Constraint:

Each page table needs to fit into a physical memory page!



Mapping of pages

Why? A page table needs consecutive space.

Memory allocated to a process

is a sparse set of nonconsecutive pages

Maximum size = # entry * page size


One-level page table cannot handle large space

- **Example:**
 - 32-bit address space with 4KB per page.
 - Page table would contain $2^{32} / 2^{12} = 1$ million entries.
 - 4 bytes per entry
- **Need a 4MB page table with contiguous space.**
 - Is there 4MB contiguous space for each process?
- **Maximum size with 4KB per page**
 - #entry = $4\text{KB} / 4\text{B} = 1\text{K}$.
 - Maximum logical space = $1\text{K} * 4\text{KB} = 4\text{MB}$.

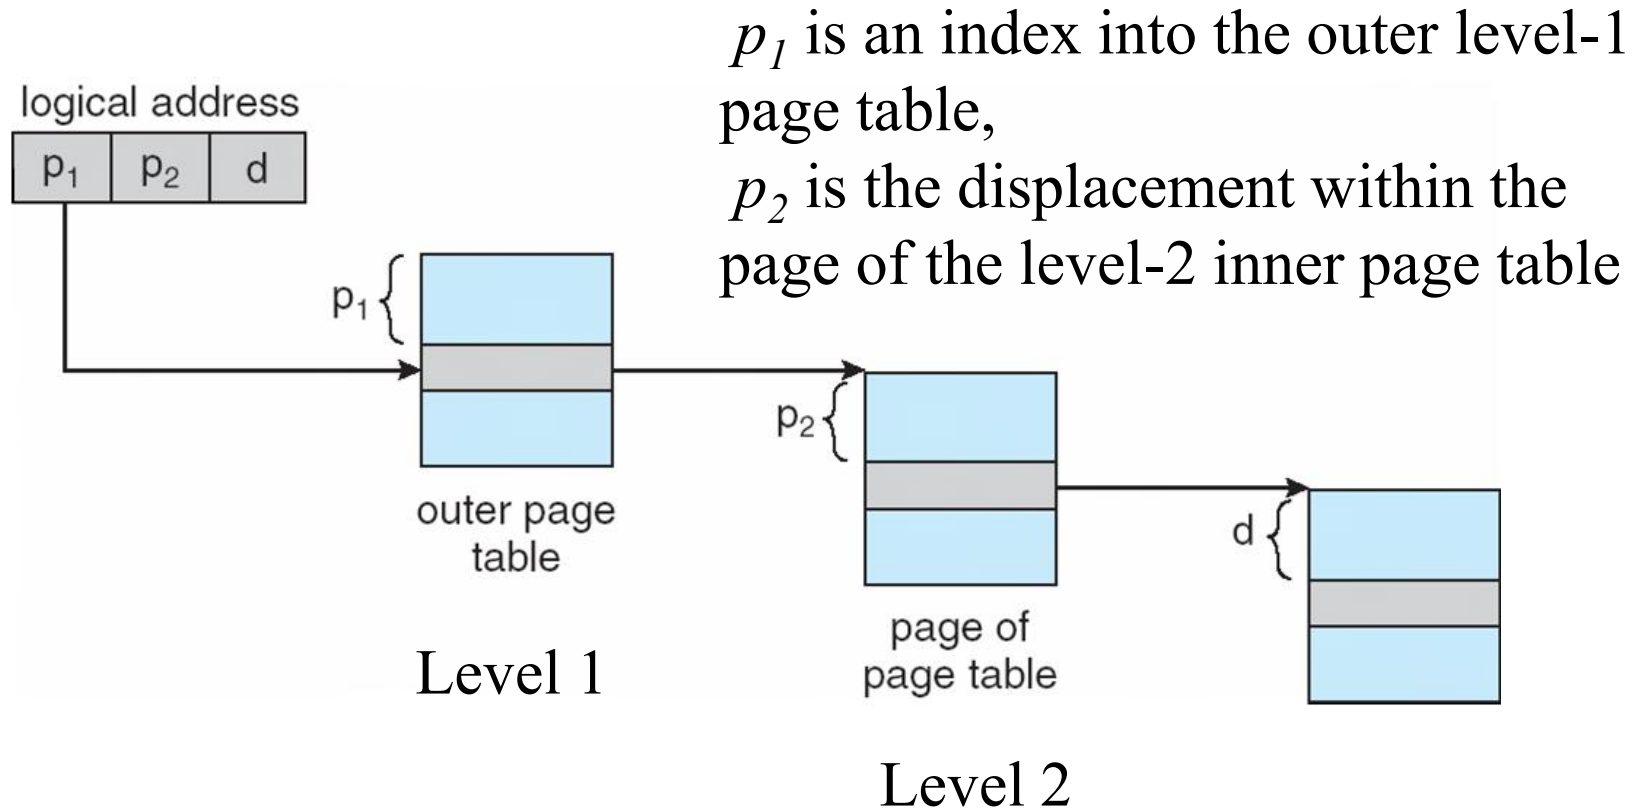
One-level Page Table: Advantage/Disadvantage

- **Pros**
 - Simple memory allocation
 - Easy to Share
- **Con: What if address space is sparse?**
 - E.g. on UNIX, code starts at 0, stack starts at $(2^{31}-1)$.
 - Cannot handle a large virtual address space
- **Con: What if table really big?**
 - Not all pages used all the time \Rightarrow would be nice to have working set of page table in memory
- **How about combining paging and segmentation?**
 - Segments with pages inside them?
 - Need some sort of multi-level translation

Two-Level Page-Table Scheme

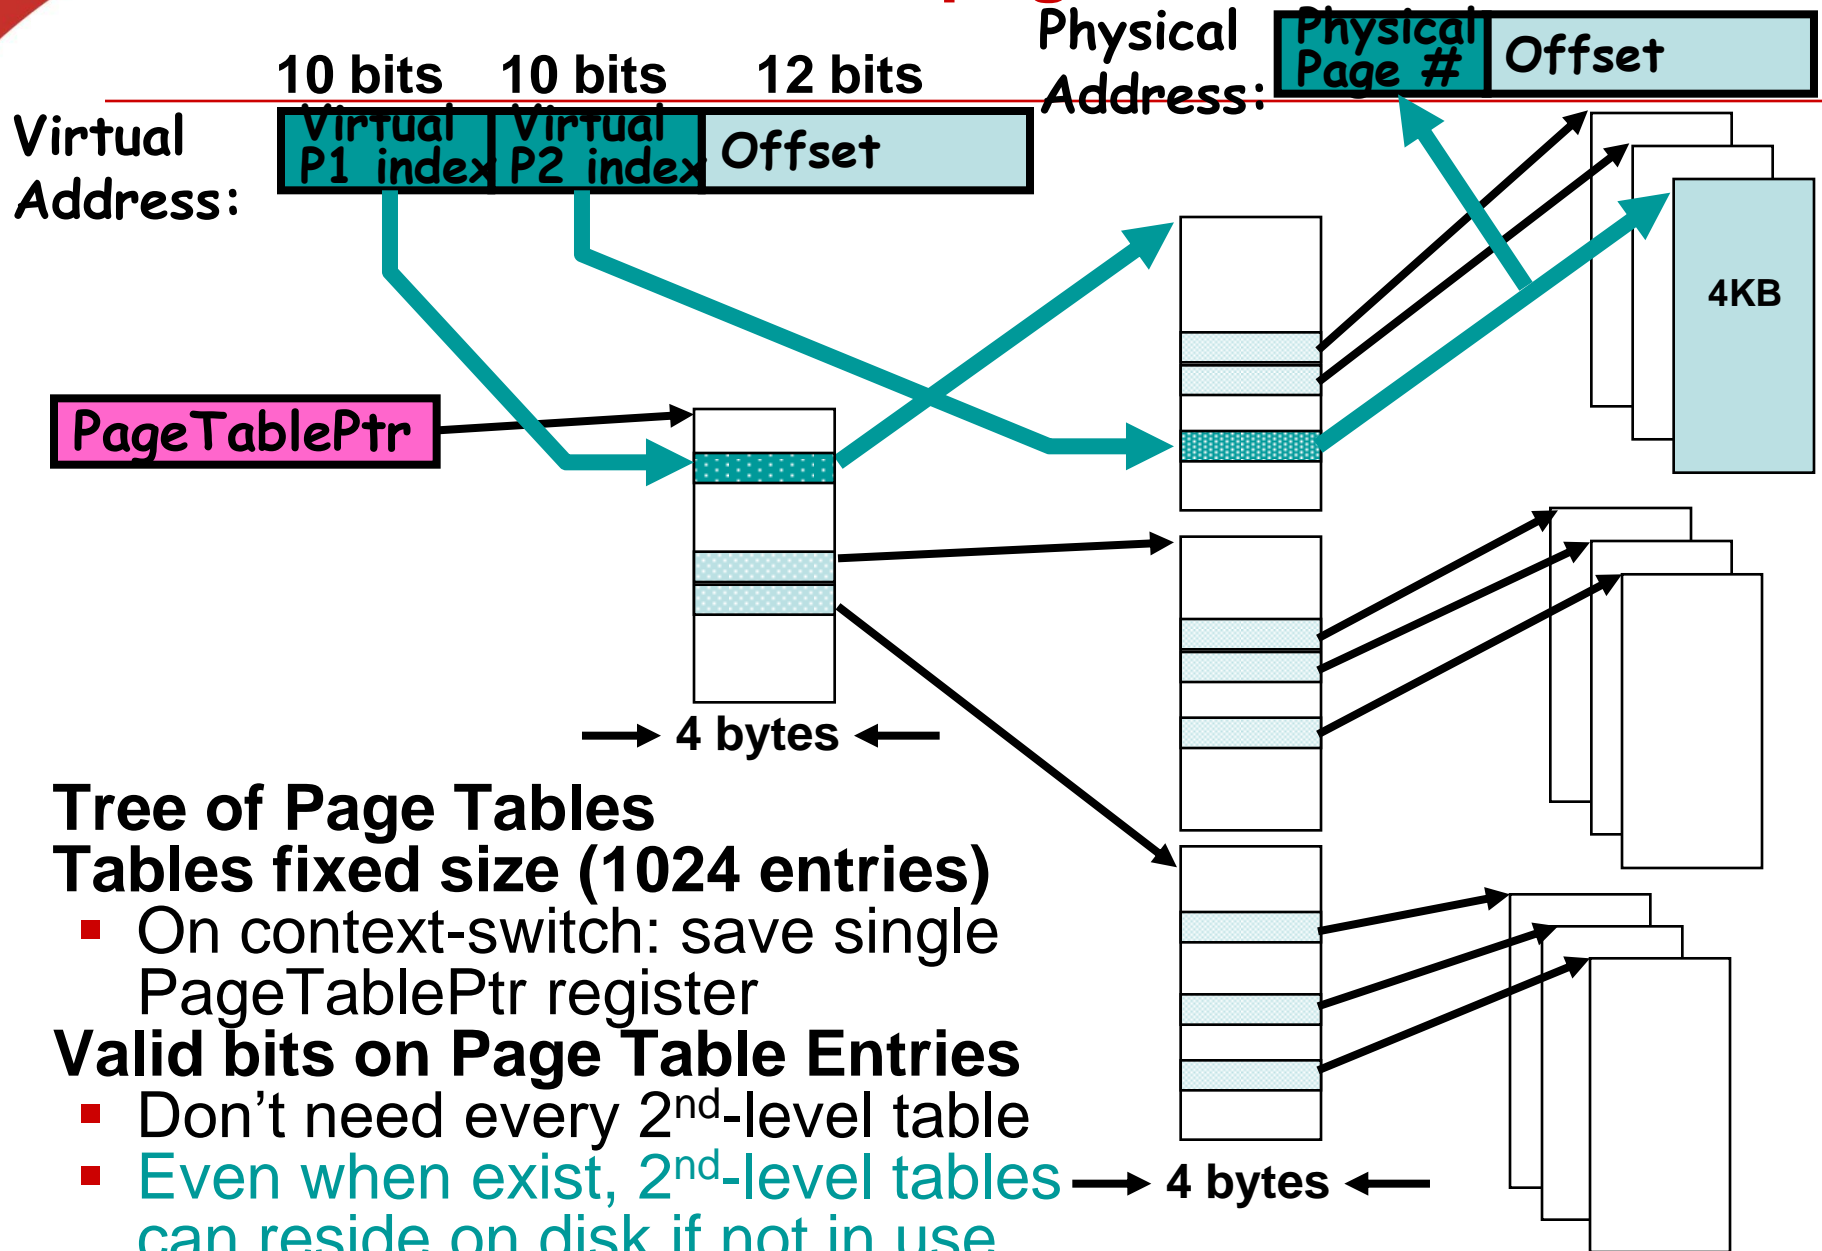


Address-Translation Scheme

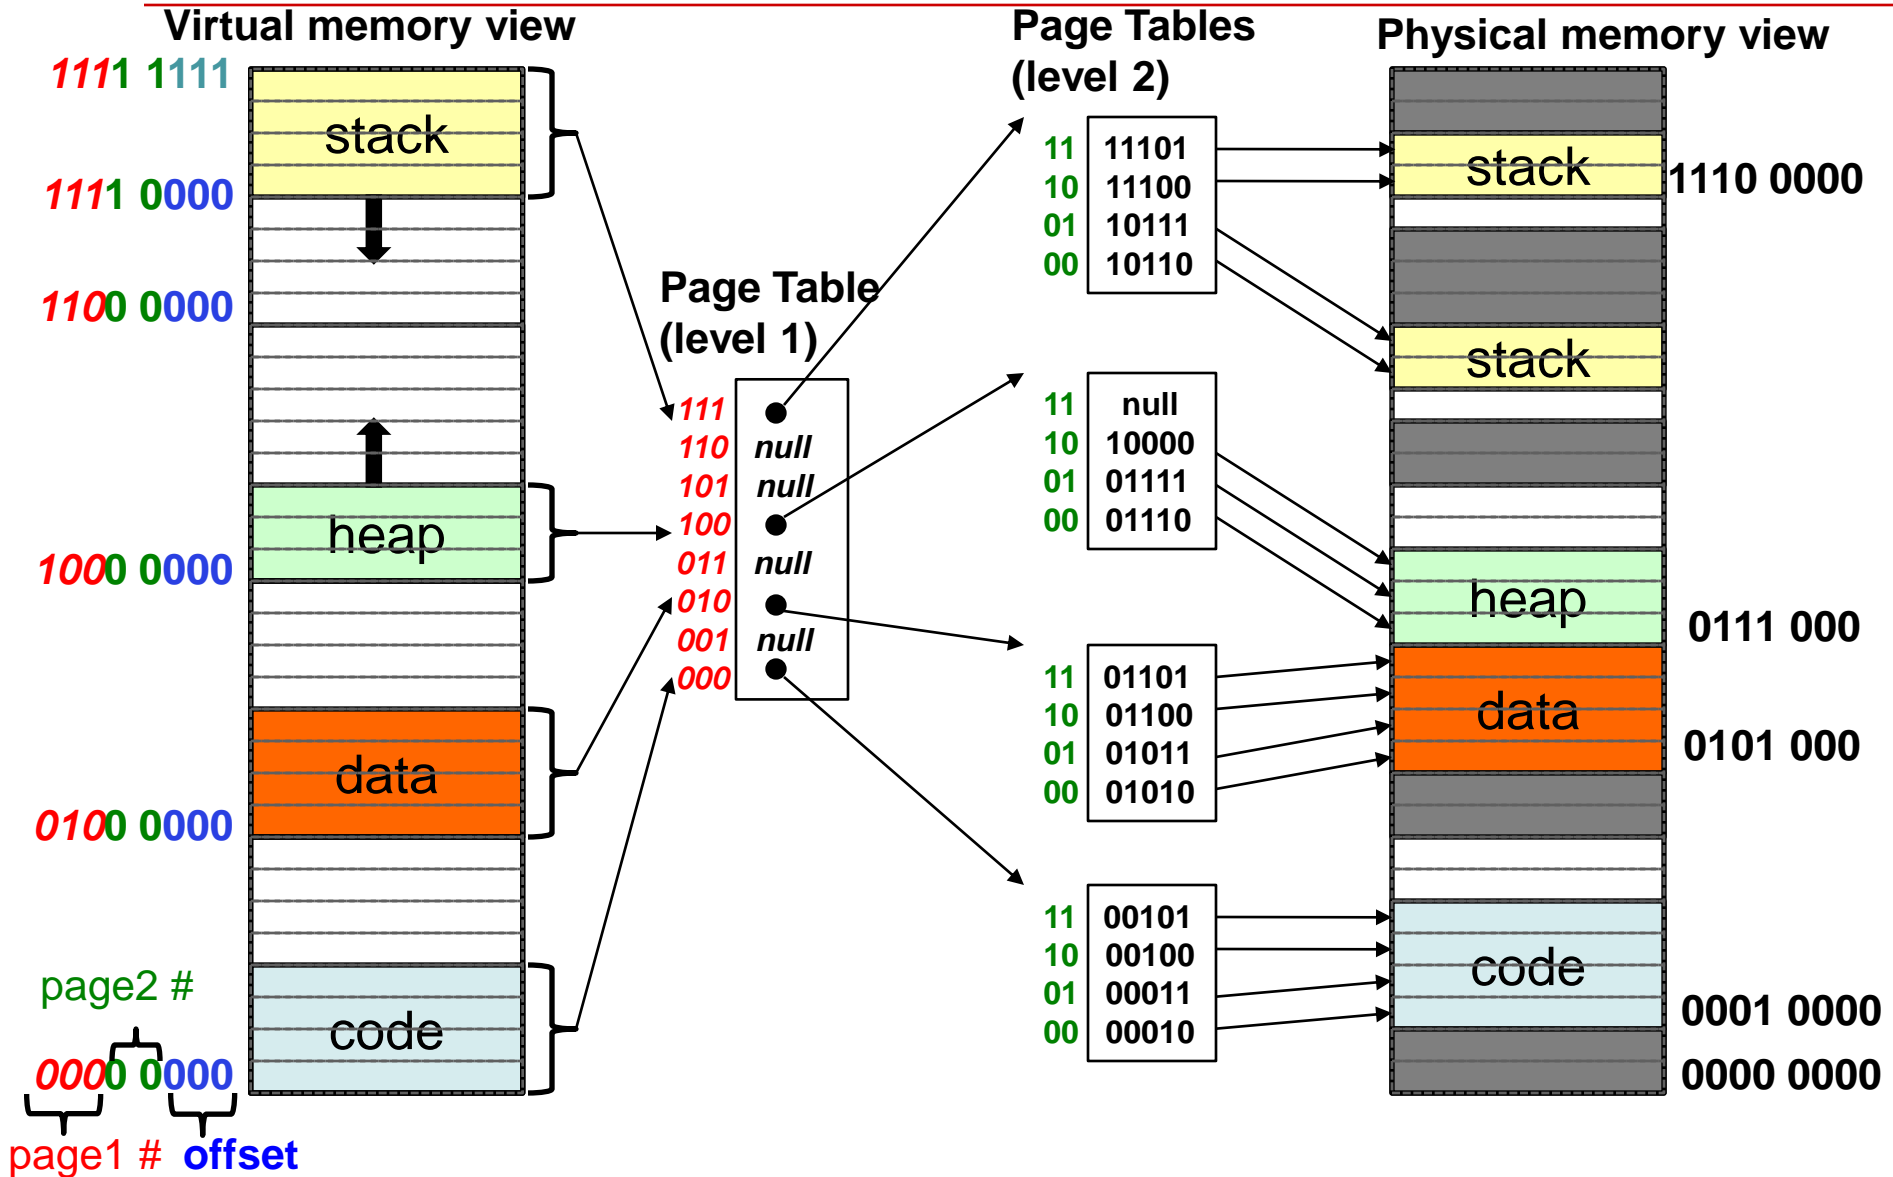


Level-2 page table gives the final physical page ID

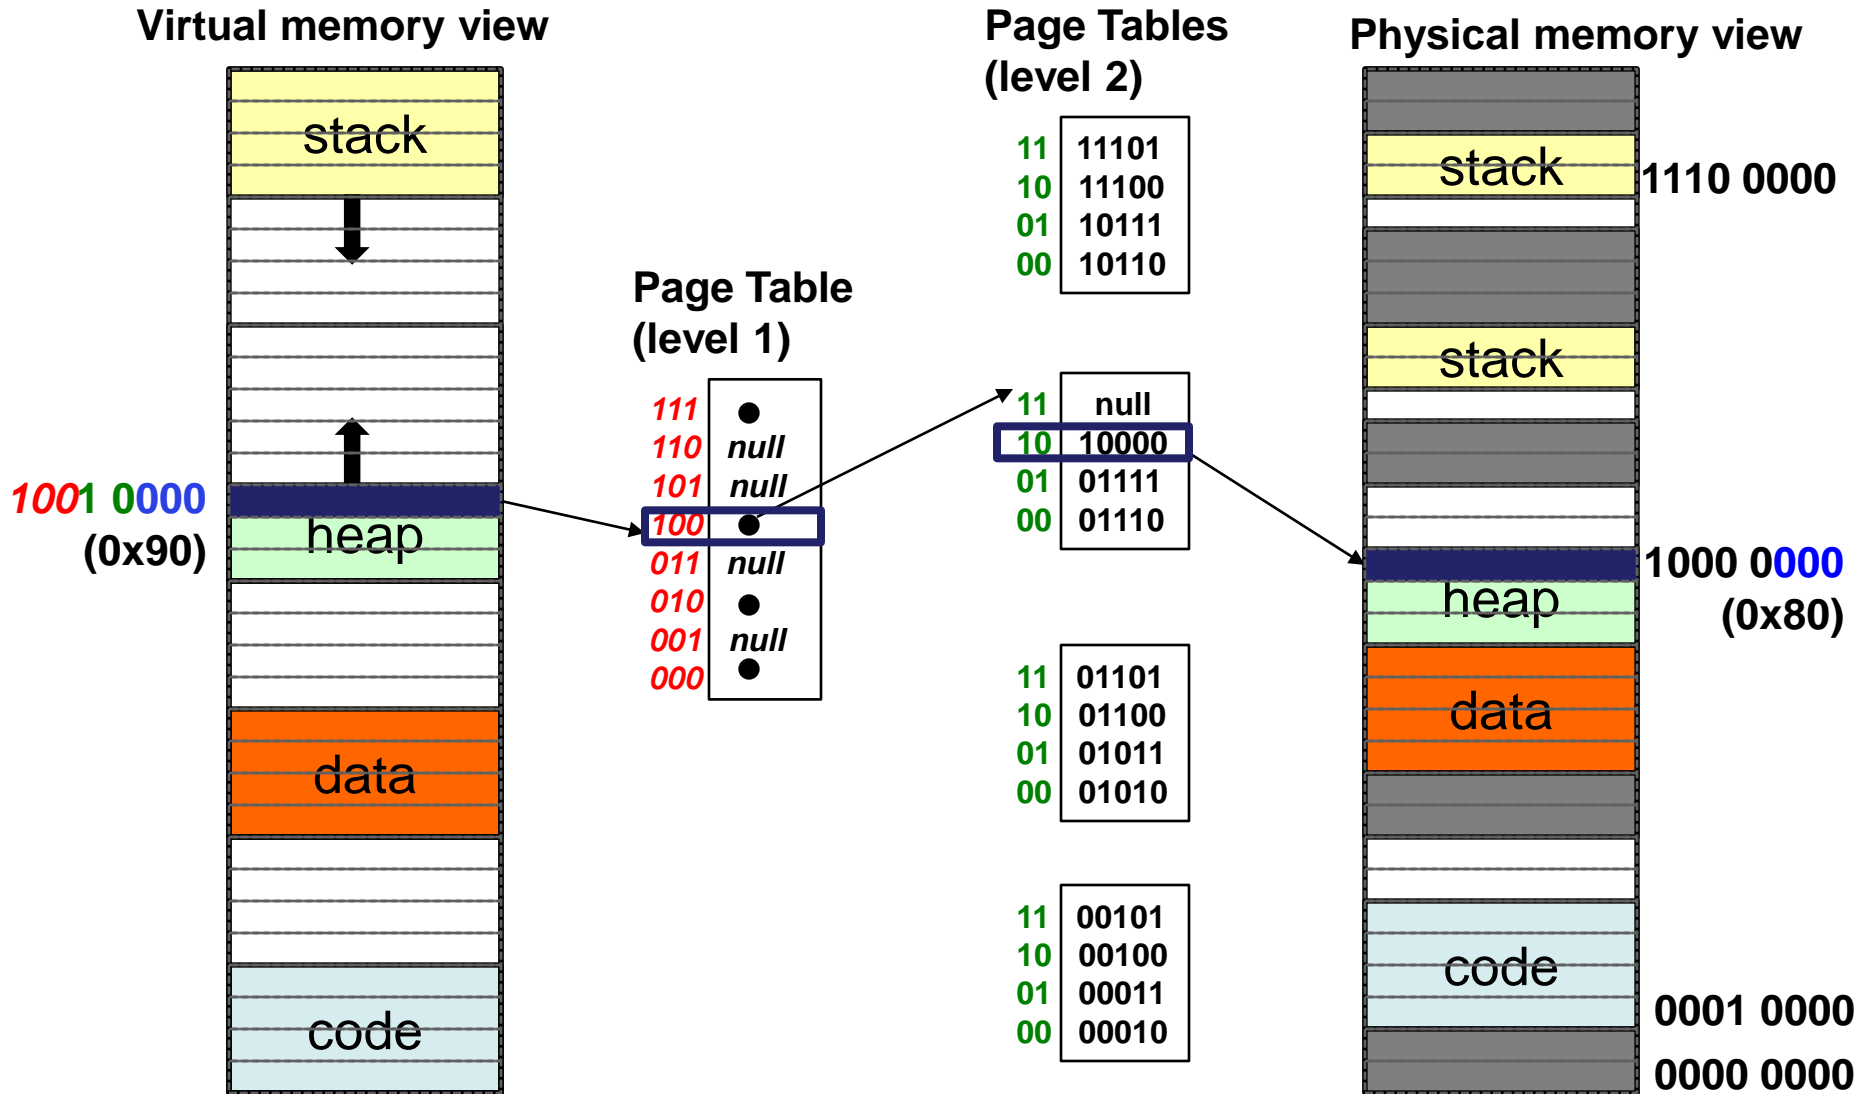
Flow of the two-level page table



Summary: Two-Level Paging

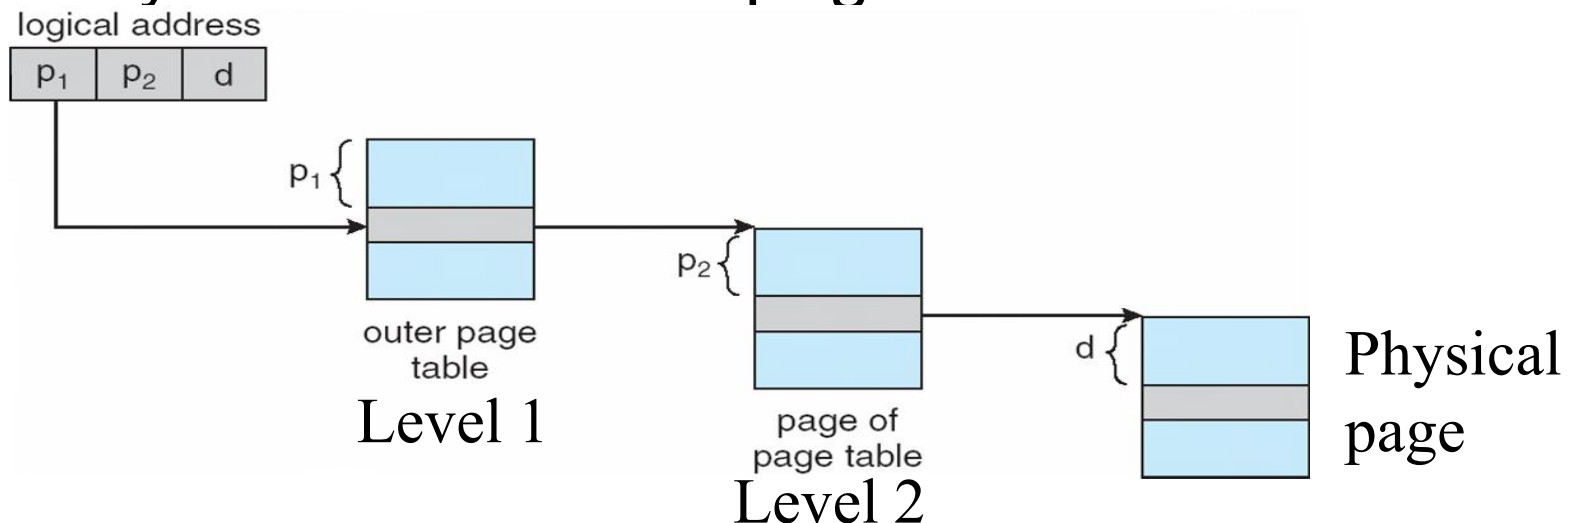


Summary: Two-Level Paging



Constraint of paging

- Bits of $d = \log(\text{page size})$
- Bits of $p_1 \geq \log(\text{\# entries in level-1 table})$
- Bits of $p_2 \geq \log(\text{\# entries in level-2 table})$
- Physical page number is limited by entry size of level-2 table.
- logical space size = $\text{\# entry in level-1 table} * \text{\# entry in level 2 table} * \text{page size}$



Analysis of a Two-Level Paging Example

- **A logical address (on 32-bit machine with 4K page size) is divided into:**
 - a page number consisting of 20 bits
 - a page offset consisting of 12 bits
- **Each entry uses 4 bytes**
- **How to build a two-level paging scheme?**
 - How many entries can a single-page table hold?
 - What are p_1 , p_2 ?

page number		page offset
p_1	p_2	d
?	?	12

Analysis of a Two-Level Paging Example

- A 1-page table with 4KB contains 1K entries and each uses 4B.
- 1K entries require 10 bits for P_1 and P_2 offset
- The page number is further divided into:
 - a 10-bit level-1 index
 - a 10-bit level-2 index

page number		page offset
p_i	p_2	d
10	10	12

What if we use 2 bytes for each table entry?

- Increased logical space size?
- Increased physical space size?

Example of maximum logical space size


- Maximum logical space size
entry in level-1 page table * # entry in level-2
page table * page size

$$= 1K * 1K * 4KB$$



$$= 2^{32} \text{ bytes}$$

$$= 4GB$$


page number		page offset
p_i	p_2	d
10	10	12



An example of three-level paging in a 64-bit address space

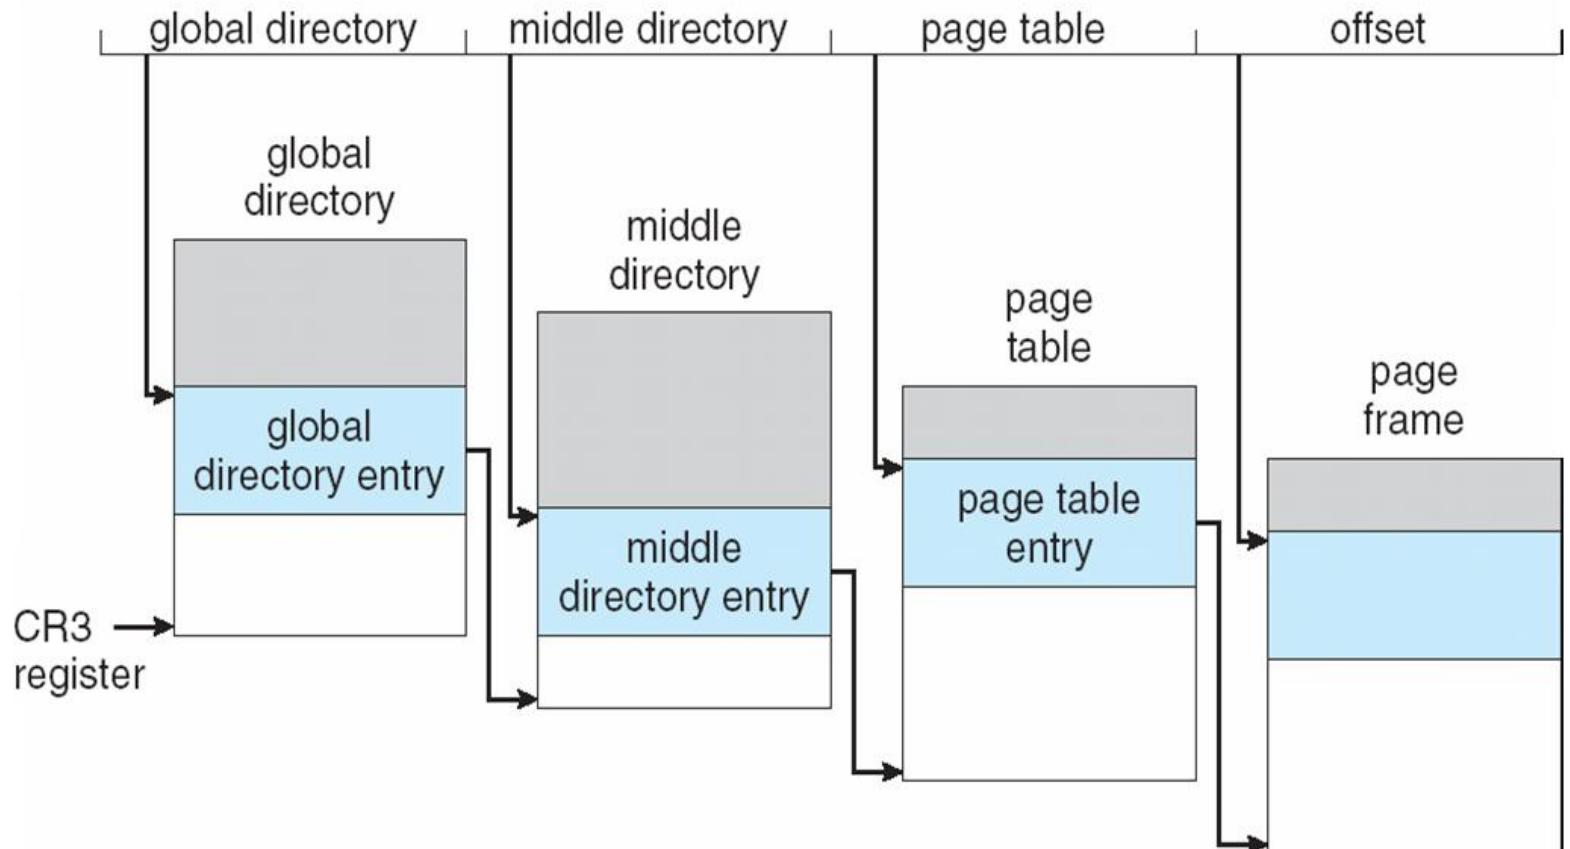



What is the maximum logical and physical space size?



Three-level Paging in Linux


Address is divided
into 4 parts



Hashed Page Tables

- **Common in address spaces > 32 bits**
- Size of page table grows proportionally as large as amount of virtual memory allocated to processes
- **Use hash table to limit the cost of search**
 - to one — or at most a few — page-table entries
 - One hash table per process
 - This page table contains a chain of elements hashing to the same location
- **Use this hash table to find the physical page of each logical page**
 - If a match is found, the corresponding physical frame is extracted

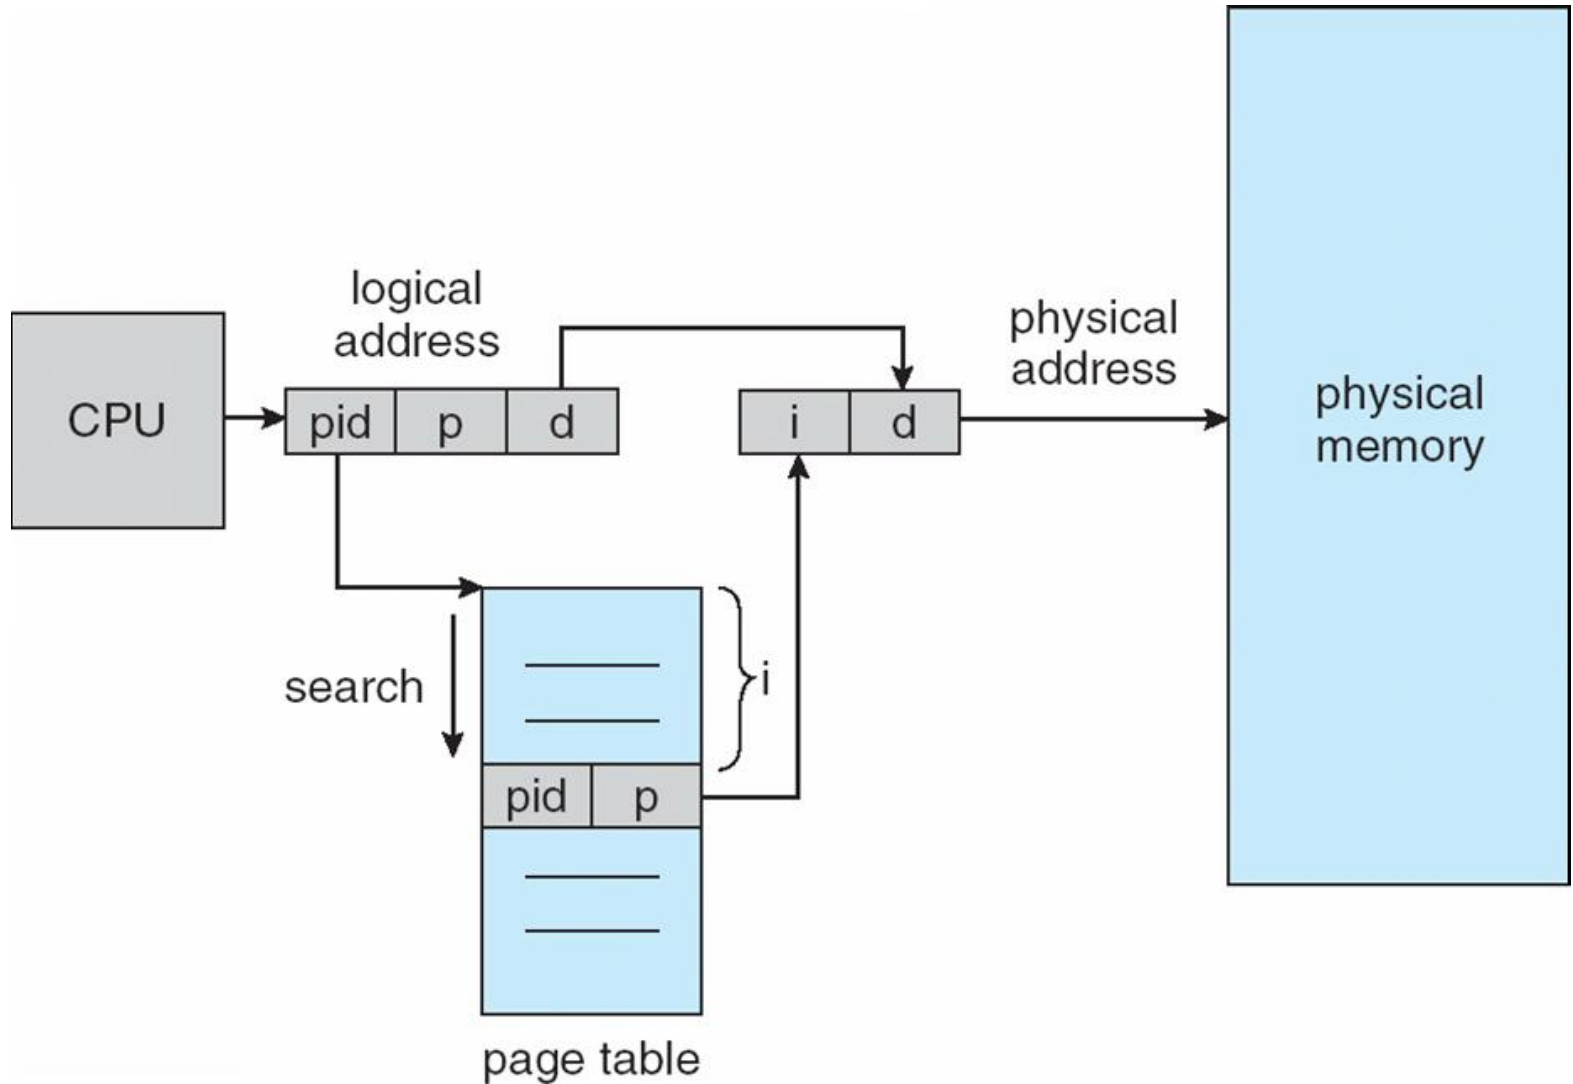
Hashed Page Table



Inverted Page Table

- **One hash table for all processes**
 - One entry for each real page of memory
 - Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page
- **Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs**

Inverted Page Table Architecture



Address Translation Comparison

	Advantages	Disadvantages
Segmentation	Fast context switching: Segment mapping maintained by CPU	External fragmentation
Paging (single-level page)	No external fragmentation, fast easy allocation	Large table size ~ virtual memory Internal fragmentation
Paged segmentation	Table size ~ # of pages in virtual memory , fast easy allocation	Multiple memory references per page access
Two-level pages		
Inverted Table	Table size ~ # of pages in physical memory	Hash function more complex

Summary

- **Page Tables**
 - Memory divided into fixed-sized chunks of memory
 - Virtual page number from virtual address mapped through page table to physical page number
 - Offset of virtual address same as physical address
 - Large page tables can be placed into virtual memory
- **Usage of page table entries**
 - Page sharing.
 - Copy on write
 - Pages on demand
 - Zero fill on demand
- **Multi-Level Tables**
 - Virtual address mapped to series of tables
 - Permit sparse population of address space
- **Inverted page table**
 - Size of page table related to physical memory size